

Part 1: EDA

Insert cells as needed below to write a short EDA/data section that summarizes the data for someone who has never opened it before.

- Answer essential questions about the dataset (observation units, time period, sample size, many of the questions above)
- Note any issues you have with the data (variable X has problem Y that needs to get addressed before using it in regressions or a prediction model because Z)
- Present any visual results you think are interesting or important

```
In [2]: import pandas as pd
import numpy as np

df = pd.read_csv('E:\\FIN377\\asgn-06-Shanshan417\\input_data2\\housing_train.csv')
```

```
In [3]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1941 entries, 0 to 1940

Data columns (total 81 columns):

#	Column	Non-Null Count	Dtype
0	parcel	1941 non-null	object
1	v_MS_SubClass	1941 non-null	int64
2	v_MS_Zoning	1941 non-null	object
3	v_Lot_Frontage	1620 non-null	float64
4	v_Lot_Area	1941 non-null	int64
5	v_Street	1941 non-null	object
6	v_Alley	136 non-null	object
7	v_Lot_Shape	1941 non-null	object
8	v_Land_Contour	1941 non-null	object
9	v_Uilities	1941 non-null	object
10	v_Lot_Config	1941 non-null	object
11	v_Land_Slope	1941 non-null	object
12	v_Neighborhood	1941 non-null	object
13	v_Condition_1	1941 non-null	object
14	v_Condition_2	1941 non-null	object
15	v_Bldg_Type	1941 non-null	object
16	v_House_Style	1941 non-null	object
17	v_Overall_Qual	1941 non-null	int64
18	v_Overall_Cond	1941 non-null	int64
19	v_Year_Built	1941 non-null	int64
20	v_Year_Remod/Add	1941 non-null	int64
21	v_Roof_Style	1941 non-null	object
22	v_Roof_Mat1	1941 non-null	object
23	v_Exterior_1st	1941 non-null	object
24	v_Exterior_2nd	1941 non-null	object
25	v_Mas_Vnr_Type	769 non-null	object
26	v_Mas_Vnr_Area	1923 non-null	float64
27	v_Exter_Qual	1941 non-null	object
28	v_Exter_Cond	1941 non-null	object
29	v_Foundation	1941 non-null	object
30	v_Bsmt_Qual	1891 non-null	object
31	v_Bsmt_Cond	1891 non-null	object
32	v_Bsmt_Exposure	1889 non-null	object
33	v_BsmtFin_Type_1	1891 non-null	object
34	v_BsmtFin_SF_1	1940 non-null	float64
35	v_BsmtFin_Type_2	1891 non-null	object
36	v_BsmtFin_SF_2	1940 non-null	float64
37	v_Bsmt_Unf_SF	1940 non-null	float64
38	v_Total_Bsmt_SF	1940 non-null	float64
39	v_Heating	1941 non-null	object
40	v_Heating_QC	1941 non-null	object
41	v_Central_Air	1941 non-null	object
42	v_Electrical	1940 non-null	object
43	v_1st_Flr_SF	1941 non-null	int64
44	v_2nd_Flr_SF	1941 non-null	int64
45	v_Low_Qual_Fin_SF	1941 non-null	int64
46	v_Gr_Liv_Area	1941 non-null	int64
47	v_Bsmt_Full_Bath	1939 non-null	float64
48	v_Bsmt_Half_Bath	1939 non-null	float64
49	v_Full_Bath	1941 non-null	int64
50	v_Half_Bath	1941 non-null	int64

```

51 v_Bedroom_AbvGr 1941 non-null int64
52 v_Kitchen_AbvGr 1941 non-null int64
53 v_Kitchen_Qual 1941 non-null object
54 v_TotRms_AbvGrd 1941 non-null int64
55 v_Functional 1941 non-null object
56 v_Fireplaces 1941 non-null int64
57 v_Fireplace_Qu 1001 non-null object
58 v_Garage_Type 1836 non-null object
59 v_Garage_Yr_Blt 1834 non-null float64
60 v_Garage_Finish 1834 non-null object
61 v_Garage_Cars 1940 non-null float64
62 v_Garage_Area 1940 non-null float64
63 v_Garage_Qual 1834 non-null object
64 v_Garage_Cond 1834 non-null object
65 v_Paved_Drive 1941 non-null object
66 v_Wood_Deck_SF 1941 non-null int64
67 v_Open_Porch_SF 1941 non-null int64
68 v_Enclosed_Porch 1941 non-null int64
69 v_3Ssn_Porch 1941 non-null int64
70 v_Screen_Porch 1941 non-null int64
71 v_Pool_Area 1941 non-null int64
72 v_Pool_QC 13 non-null object
73 v_Fence 365 non-null object
74 v_Misc_Feature 63 non-null object
75 v_Misc_Val 1941 non-null int64
76 v_Mo_Sold 1941 non-null int64
77 v_Yr_Sold 1941 non-null int64
78 v_Sale_Type 1941 non-null object
79 v_Sale_Condition 1941 non-null object
80 v_SalePrice 1941 non-null int64

```

dtypes: float64(11), int64(26), object(44)

memory usage: 1.2+ MB

In [4]: `df.describe()`

```

Out[4]:
      v_MS_SubClass  v_Lot_Frontage  v_Lot_Area  v_Overall_Qual  v_Overall_Cond
count      1941.000000      1620.000000      1941.000000      1941.000000      1941.000000
mean         58.088614         69.301235      10284.770222         6.113344         5.568264
std         42.946015         23.978101       7832.295527         1.401594         1.087465
min          20.000000         21.000000      1470.000000         1.000000         1.000000
25%          20.000000         58.000000      7420.000000         5.000000         5.000000
50%          50.000000         68.000000      9450.000000         6.000000         5.000000
75%          70.000000         80.000000     11631.000000         7.000000         6.000000
max         190.000000        313.000000     164660.000000        10.000000         9.000000

```

8 rows × 37 columns

In [5]: `df.isnull().sum()`

```
Out[5]: parcel          0
v_MS_SubClass          0
v_MS_Zoning            0
v_Lot_Frontage        321
v_Lot_Area             0
...
v_Mo_Sold              0
v_Yr_Sold              0
v_Sale_Type            0
v_Sale_Condition       0
v_SalePrice            0
Length: 81, dtype: int64
```

```
In [6]: df.nunique().sort_values()
```

```
Out[6]: v_Central_Air      2
v_Street                  2
v_Alley                   2
v_Uilities                 2
v_Yr_Sold                  3
...
v_1st_Flr_SF              901
v_Bsmt_Unf_SF             938
v_Gr_Liv_Area             1045
v_Lot_Area                1413
parcel                    1941
Length: 81, dtype: int64
```

```
In [7]: missing = df.isnull().mean().sort_values(ascending=False)
print(missing[missing > 0])
```

```

v_Pool_QC            0.993302
v_Misc_Feature       0.967543
v_Alley              0.929933
v_Fence              0.811953
v_Mas_Vnr_Type       0.603812
v_Fireplace_Qu       0.484286
v_Lot_Frontage       0.165379
v_Garage_Cond        0.055126
v_Garage_Finish      0.055126
v_Garage_Yr_Blt      0.055126
v_Garage_Qual        0.055126
v_Garage_Type        0.054096
v_Bsmt_Exposure      0.026790
v_Bsmt_Qual          0.025760
v_Bsmt_Cond          0.025760
v_BsmtFin_Type_1     0.025760
v_BsmtFin_Type_2     0.025760
v_Mas_Vnr_Area       0.009274
v_Bsmt_Half_Bath     0.001030
v_Bsmt_Full_Bath     0.001030
v_BsmtFin_SF_1       0.000515
v_Garage_Cars        0.000515
v_Electrical         0.000515
v_Total_Bsmt_SF      0.000515
v_Bsmt_Unf_SF        0.000515
v_BsmtFin_SF_2       0.000515
v_Garage_Area        0.000515
dtype: float64

```

```

In [8]: import seaborn as sns
import matplotlib.pyplot as plt

missing_percent = df.isnull().mean().sort_values(ascending=False) * 100
missing_percent = missing_percent[missing_percent > 0]

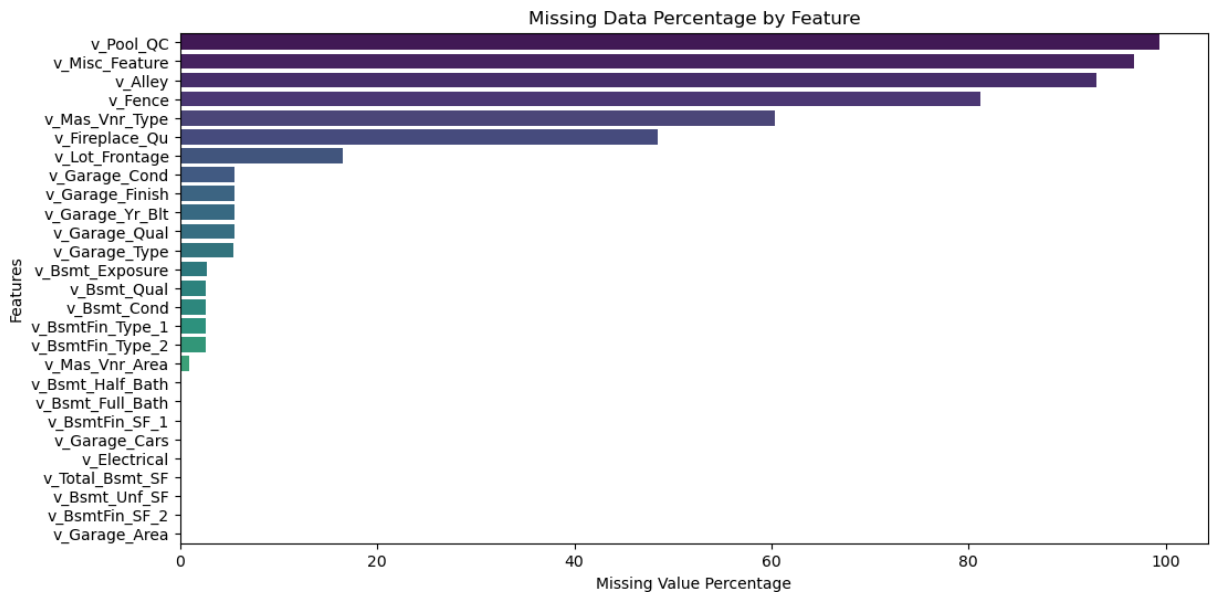
plt.figure(figsize=(12, 6))
sns.barplot(x=missing_percent.values, y=missing_percent.index, palette="viridis")
plt.xlabel("Missing Value Percentage")
plt.ylabel("Features")
plt.title("Missing Data Percentage by Feature")
plt.show()

```

C:\Users\lenovo\AppData\Local\Temp\ipykernel_52060\3277439945.py:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=missing_percent.values, y=missing_percent.index, palette="viridis")
```



```
In [9]: none_fill = [
        'v_Pool_QC', 'v_Misc_Feature', 'v_Alley', 'v_Fence',
        'v_Fireplace_Qu', 'v_Garage_Type', 'v_Garage_Finish',
        'v_Garage_Qual', 'v_Garage_Cond', 'v_Bsmt_Exposure',
        'v_Bsmt_Qual', 'v_Bsmt_Cond', 'v_BsmtFin_Type_1',
        'v_BsmtFin_Type_2', 'v_Electrical'
      ]
for col in none_fill:
    df[col] = df[col].fillna('None')
```

```
In [10]: zero_fill = [
        'v_Garage_Yr_Blt', 'v_Mas_Vnr_Area', 'v_Bsmt_Full_Bath',
        'v_Bsmt_Half_Bath', 'v_BsmtFin_SF_1', 'v_BsmtFin_SF_2',
        'v_Bsmt_Unf_SF', 'v_Total_Bsmt_SF', 'v_Garage_Cars',
        'v_Garage_Area'
      ]
for col in zero_fill:
    df[col] = df[col].fillna(0)
```

```
In [11]: mode_fill = ['v_Electrical']
for col in mode_fill:
    df[col] = df[col].fillna(df[col].mode()[0])
```

```
In [12]: from sklearn.ensemble import RandomForestRegressor

features = ['v_Overall_Qual', 'v_Lot_Area', 'v_Year_Built', 'v_Gr_Liv_Area']
lot_data = df[features + ['v_Lot_Frontage']]
train_data = lot_data[lot_data['v_Lot_Frontage'].notnull()]
predict_data = lot_data[lot_data['v_Lot_Frontage'].isnull()]
X_train = train_data[features]
y_train = train_data['v_Lot_Frontage']
model = RandomForestRegressor(random_state=0, n_estimators=100)
model.fit(X_train, y_train)
X_predict = predict_data[features]
predicted_values = model.predict(X_predict)
df.loc[df['v_Lot_Frontage'].isnull(), 'v_Lot_Frontage'] = predicted_values
```

```
In [13]: missing = df.isnull().mean().sort_values(ascending=False)
print(missing[missing > 0])
```

```
v_Mas_Vnr_Type    0.603812
dtype: float64
```

```
In [14]: df['v_Mas_Vnr_Type'] = df['v_Mas_Vnr_Type'].fillna('None')
```

```
In [15]: missing = df.isnull().mean().sort_values(ascending=False)
print(missing[missing > 0])
```

```
Series([], dtype: float64)
```


```
In [16]: numerical = df.select_dtypes(include=['int64', 'float64']).columns
categorical = df.select_dtypes(include=['object']).columns
print("Numerical:", numerical)
print("Categorical:", categorical)
```

```
Numerical: Index(['v_MS_SubClass', 'v_Lot_Frontage', 'v_Lot_Area', 'v_Overall_Qual',
                 'v_Overall_Cond', 'v_Year_Built', 'v_Year_Remod/Add', 'v_Mas_Vnr_Area',
                 'v_BsmtFin_SF_1', 'v_BsmtFin_SF_2', 'v_Bsmt_Unf_SF', 'v_Total_Bsmt_SF',
                 'v_1st_Flr_SF', 'v_2nd_Flr_SF', 'v_Low_Qual_Fin_SF', 'v_Gr_Liv_Area',
                 'v_Bsmt_Full_Bath', 'v_Bsmt_Half_Bath', 'v_Full_Bath', 'v_Half_Bath',
                 'v_Bedroom_AbvGr', 'v_Kitchen_AbvGr', 'v_TotRms_AbvGrd', 'v_Fireplaces',
                 'v_Garage_Yr_Blt', 'v_Garage_Cars', 'v_Garage_Area', 'v_Wood_Deck_SF',
                 'v_Open_Porch_SF', 'v_Enclosed_Porch', 'v_3Ssn_Porch', 'v_Screen_Porch',
                 'v_Pool_Area', 'v_Misc_Val', 'v_Mo_Sold', 'v_Yr_Sold', 'v_SalePrice'],
                 dtype='object')
```



```
Categorical: Index(['parcel', 'v_MS_Zoning', 'v_Street', 'v_Alley', 'v_Lot_Shape',
                   'v_Land_Contour', 'v_Uutilities', 'v_Lot_Config', 'v_Land_Slope',
                   'v_Neighborhood', 'v_Condition_1', 'v_Condition_2', 'v_Bldg_Type',
                   'v_House_Style', 'v_Roof_Style', 'v_Roof_Matl', 'v_Exterior_1st',
                   'v_Exterior_2nd', 'v_Mas_Vnr_Type', 'v_Exter_Qual', 'v_Exter_Cond',
                   'v_Foundation', 'v_Bsmt_Qual', 'v_Bsmt_Cond', 'v_Bsmt_Exposure',
                   'v_BsmtFin_Type_1', 'v_BsmtFin_Type_2', 'v_Heating', 'v_Heating_QC',
                   'v_Central_Air', 'v_Electrical', 'v_Kitchen_Qual', 'v_Functional',
                   'v_Fireplace_Qu', 'v_Garage_Type', 'v_Garage_Finish', 'v_Garage_Qual',
                   'v_Garage_Cond', 'v_Paved_Drive', 'v_Pool_QC', 'v_Fence',
                   'v_Misc_Feature', 'v_Sale_Type', 'v_Sale_Condition'],
                   dtype='object')
```

```
In [17]: corr = df.corr(numeric_only=True)
top_corr = corr['v_SalePrice'].abs().sort_values(ascending=False).head(11)
top_features = top_corr.index


plt.figure(figsize=(10, 8))
sns.heatmap(df[top_features].corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Top Correlated Features with v_SalePrice')
plt.show()
```




```
In [18]: for col in numerical:
          if col != 'v_SalePrice':
              sns.scatterplot(x=df[col], y=df['v_SalePrice'])
              plt.title(f'{col} vs v_SalePrice')
              plt.show()
```


v_Lot_Area vs v_SalePrice




v_Overall_Qual vs v_SalePrice




v_Overall_Cond vs v_SalePrice




v_Year_Built vs v_SalePrice

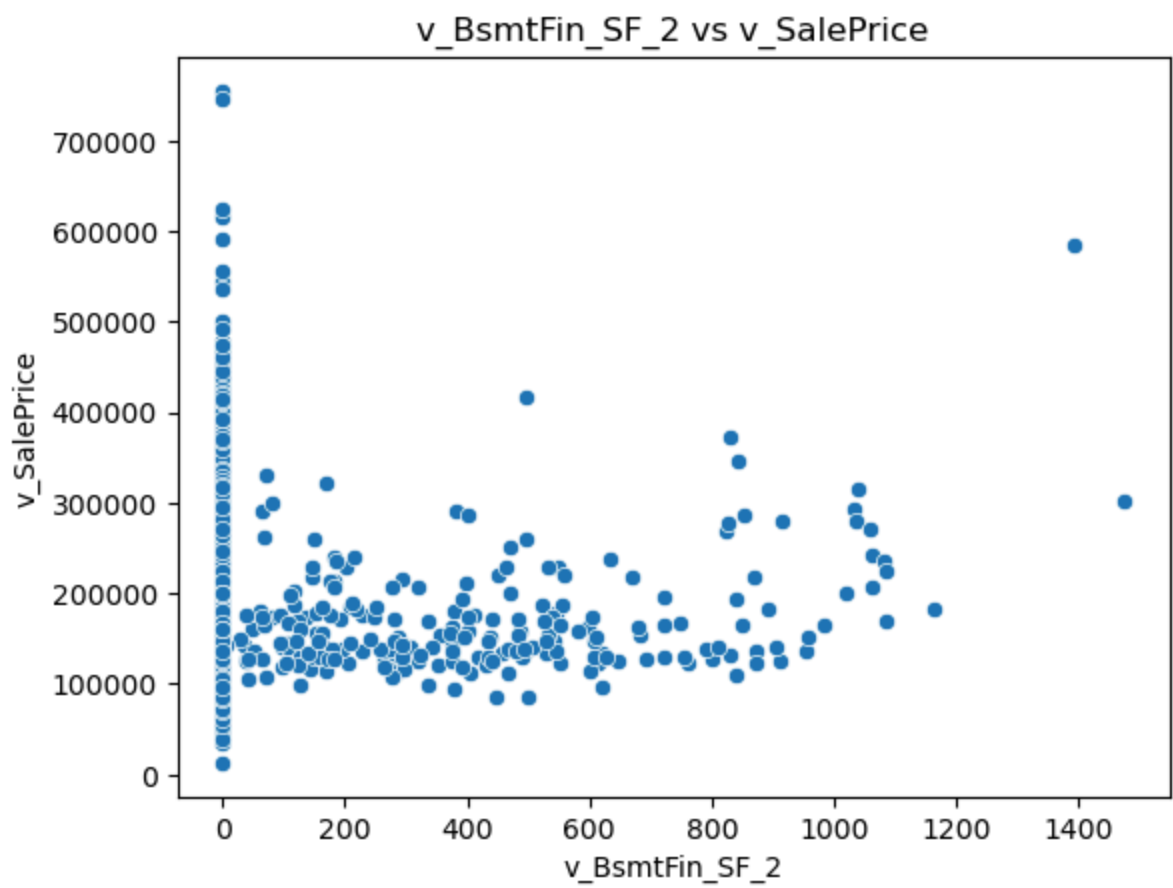
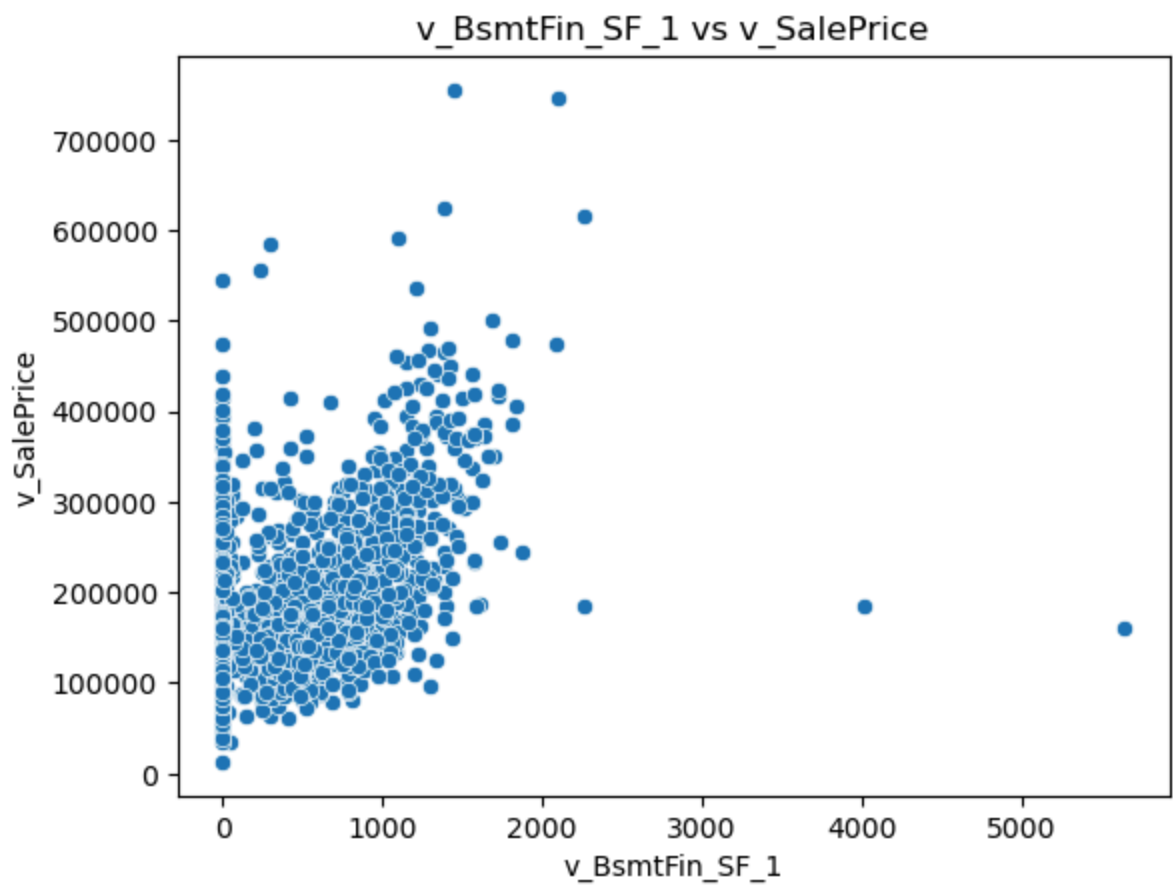




v_Year_Remod/Add vs v_SalePrice

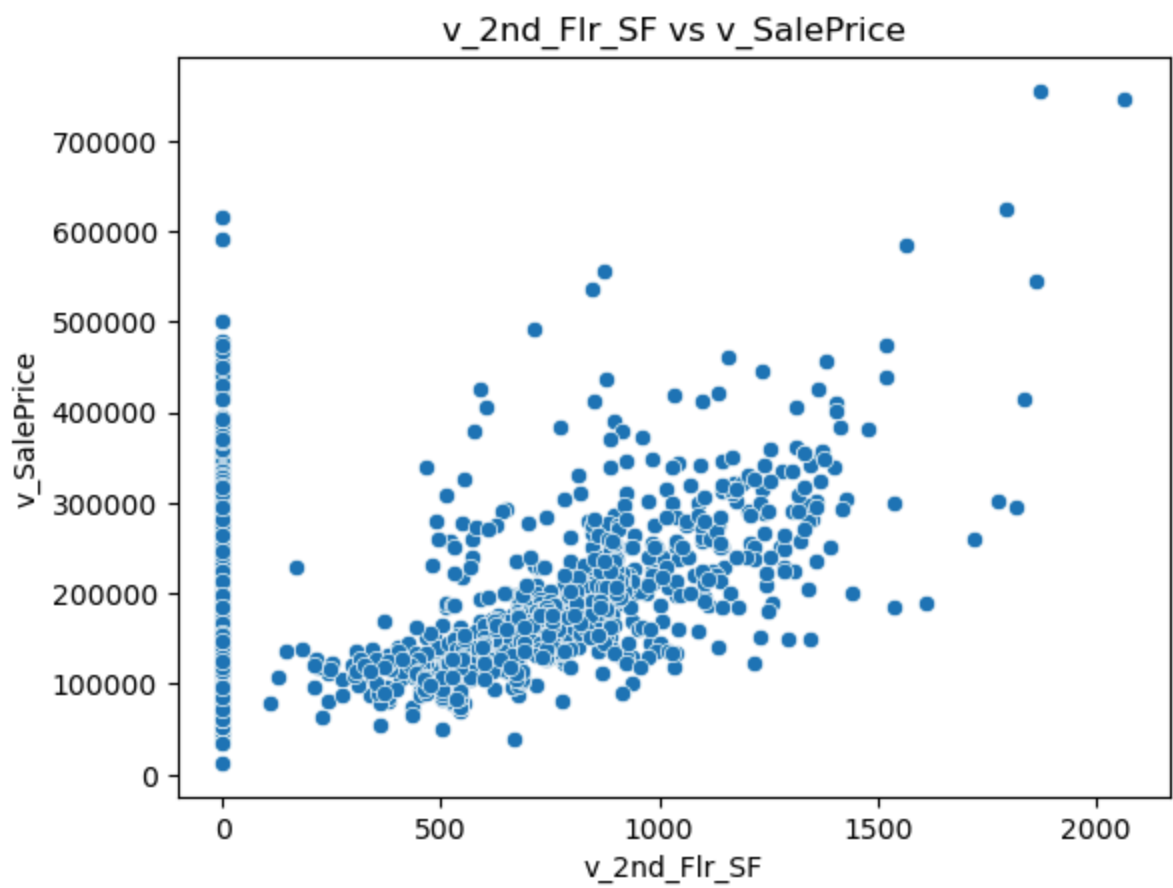




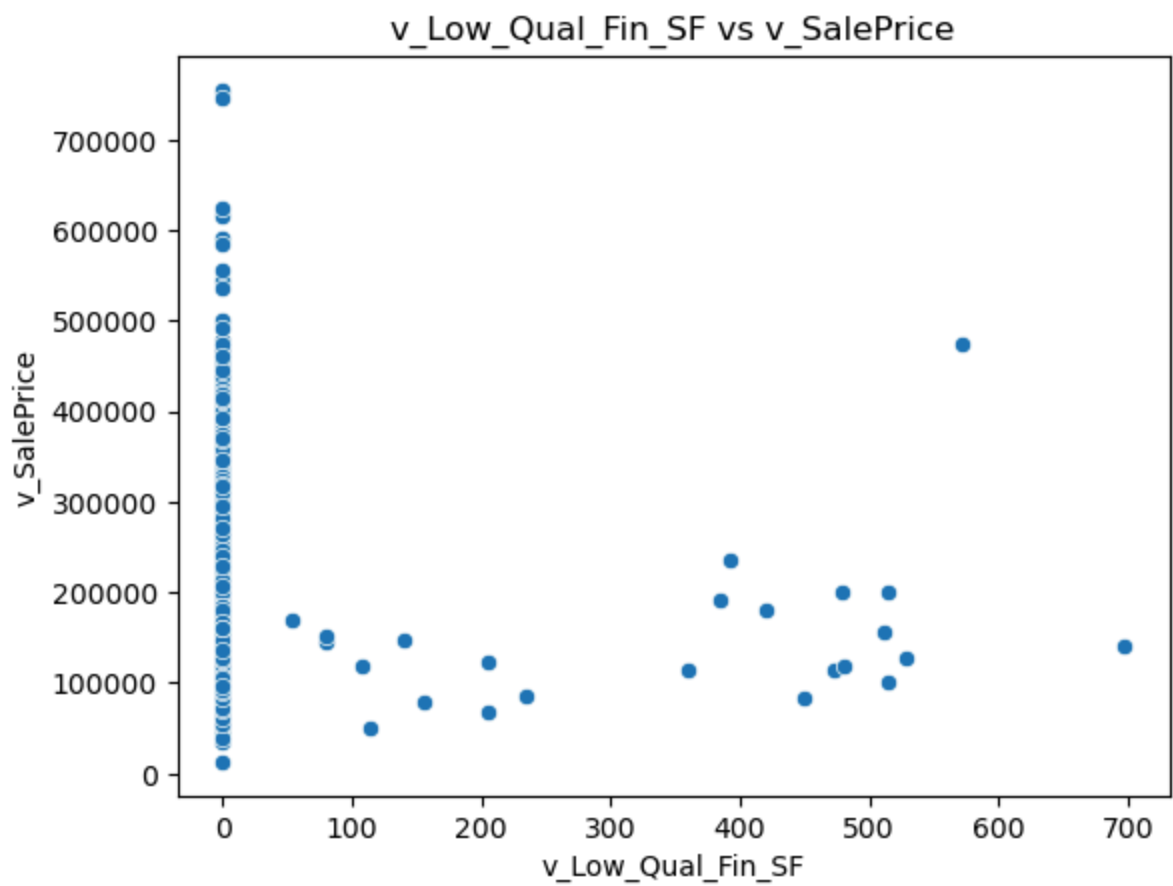
v_Mas_Vnr_Area vs v_SalePrice

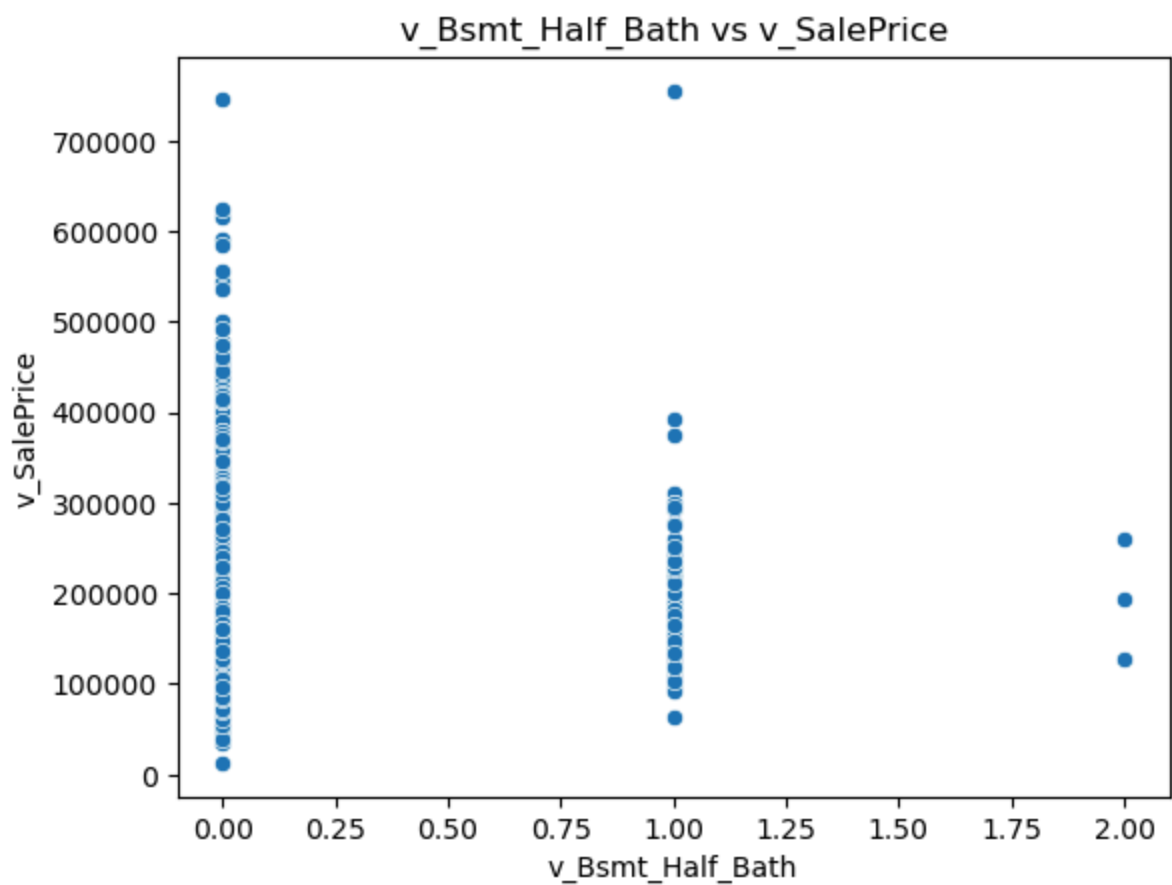
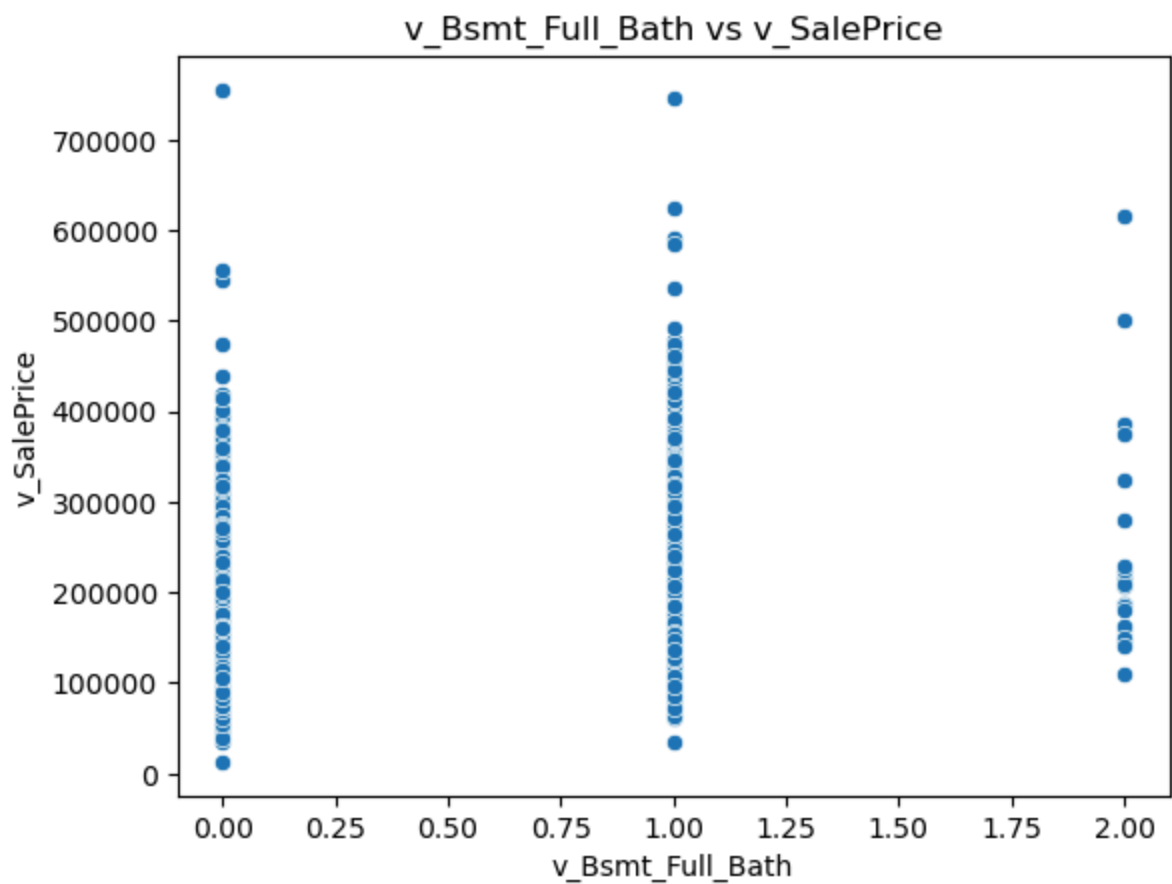















v_Full_Bath vs v_SalePrice




v_Half_Bath vs v_SalePrice

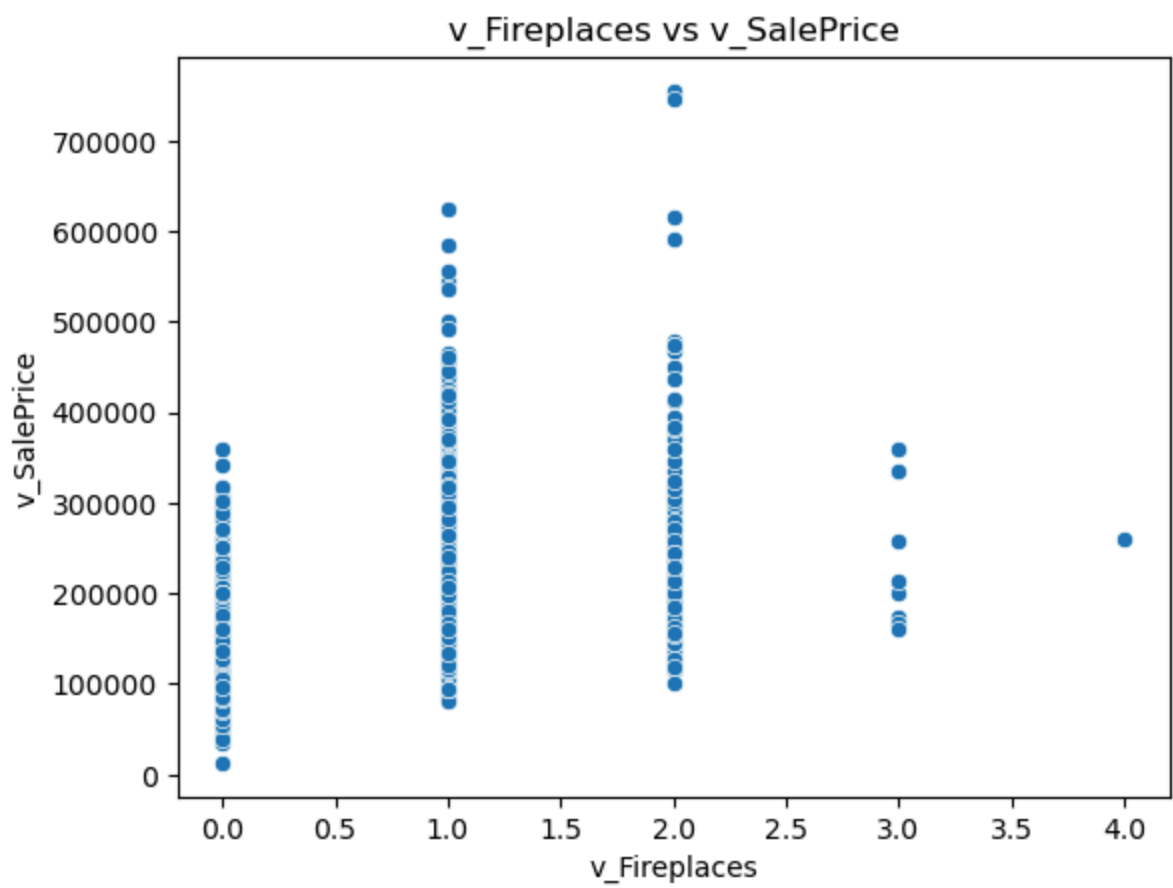
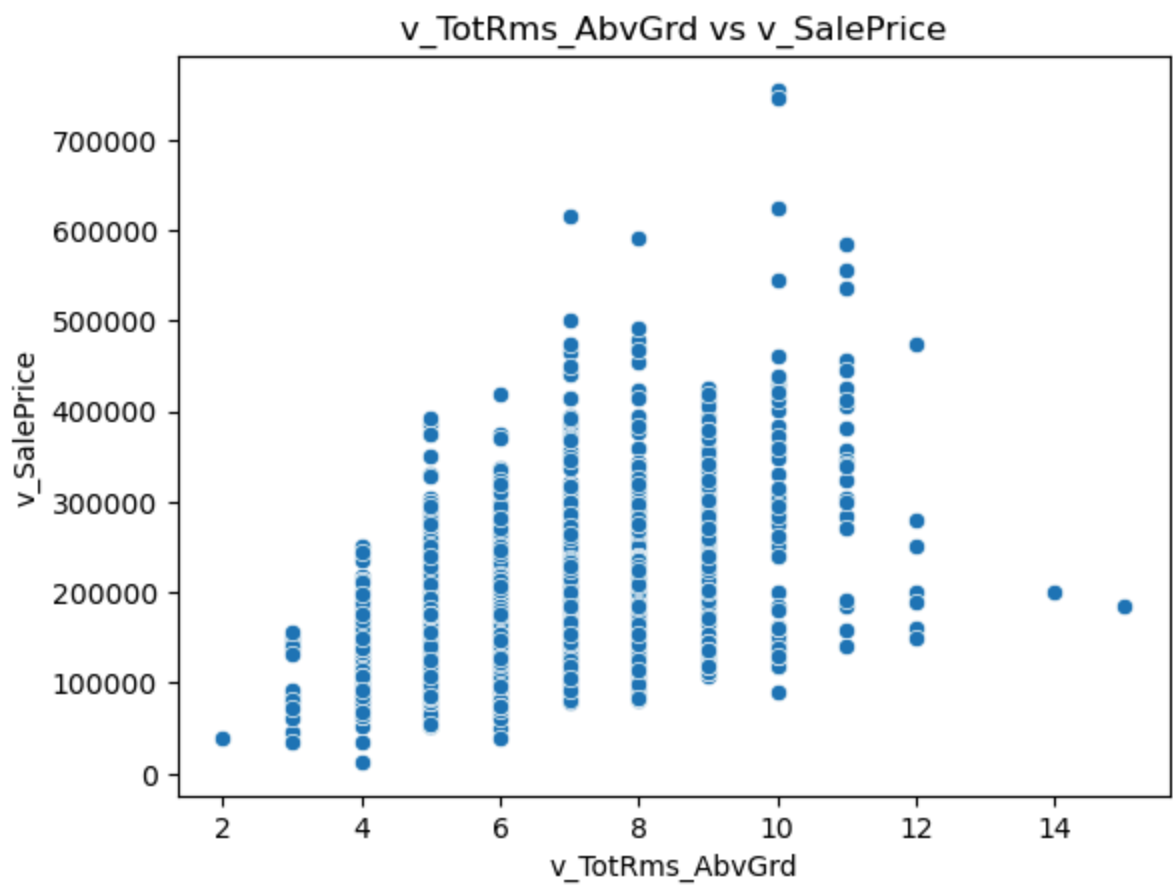


v_Bedroom_AbvGr vs v_SalePrice




v_Kitchen_AbvGr vs v_SalePrice







v_Garage_Yr_Blt vs v_SalePrice




v_Garage_Cars vs v_SalePrice

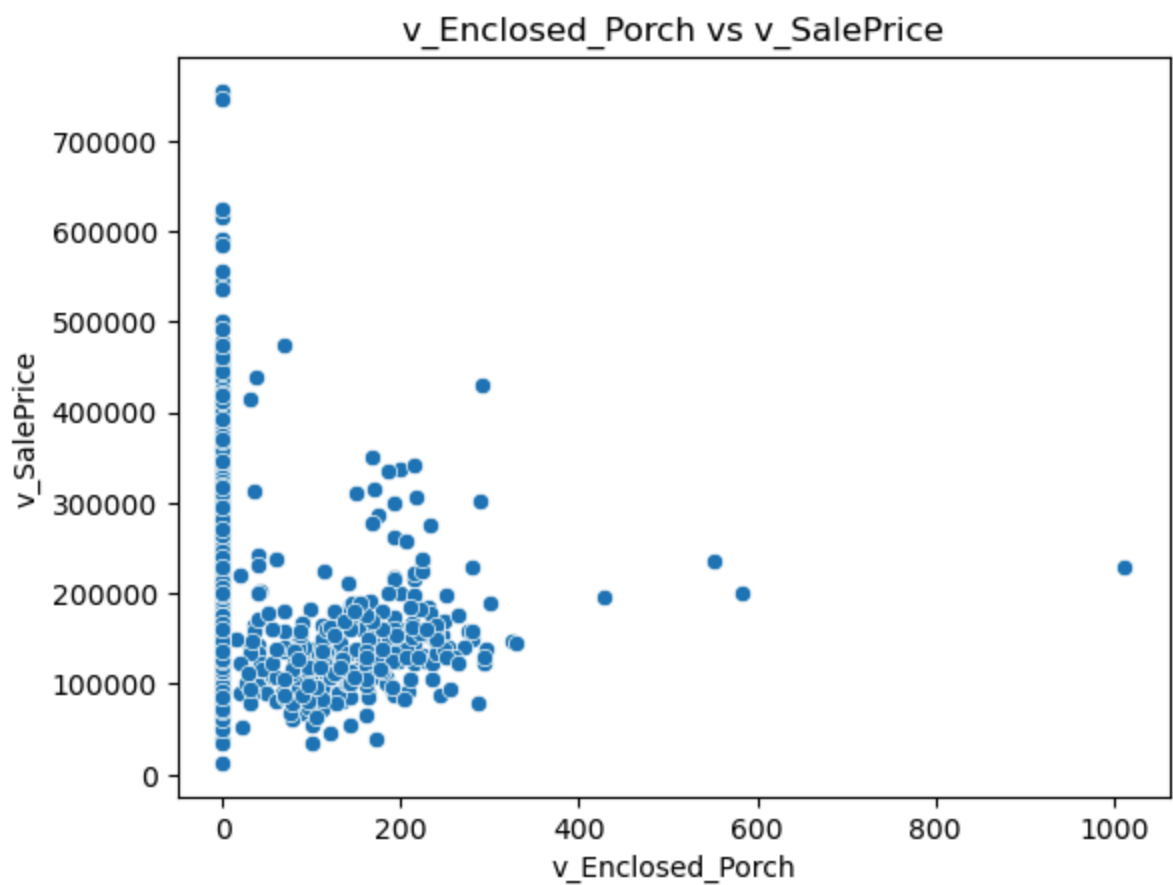



v_Garage_Area vs v_SalePrice




v_Wood_Deck_SF vs v_SalePrice







v_3Ssn_Porch vs v_SalePrice




v_Screen_Porch vs v_SalePrice




v_Pool_Area vs v_SalePrice






v_Misc_Val vs v_SalePrice








```
In [19]: for col in categorical:
          if df[col].nunique() < 10:
              sns.boxplot(x=df[col], y=df['v_SalePrice'])
              plt.title(f'{col} vs v_SalePrice')
              plt.xticks(rotation=45)
              plt.show()
```














v_Lot_Config vs v_SalePrice






v_Condition_2 vs v_SalePrice







v_House_Style vs v_SalePrice







v_Roof_Style vs v_SalePrice







v_Roof_Matl vs v_SalePrice


























v_Functional vs v_SalePrice







v_Fireplace_Qu vs v_SalePrice





v_Garage_Type vs v_SalePrice













v_Misc_Feature vs v_SalePrice





Part 2: Running Regressions

Run these regressions on the RAW data, even if you found data issues that you think should be addressed.

Insert cells as needed below to run these regressions. Note that i is indexing a given house, and t indexes the year of sale.

Note: If you are using VS Code, these might not display correctly. Add a "\" in front of the underscores in the variable names, so `\text{v_Lot_Area}` becomes

`\text{v_Lot_Area}`.

1. $\text{Sale Price}_{i,t} = \alpha + \beta_1 * \text{v_Lot_Area}$
2. $\text{Sale Price}_{i,t} = \alpha + \beta_1 * \log(\text{v_Lot_Area})$
3. $\log(\text{Sale Price}_{i,t}) = \alpha + \beta_1 * \text{v_Lot_Area}$
4. $\log(\text{Sale Price}_{i,t}) = \alpha + \beta_1 * \log(\text{v_Lot_Area})$
5. $\log(\text{Sale Price}_{i,t}) = \alpha + \beta_1 * \text{v_Yr_Sold}$
6. $\log(\text{Sale Price}_{i,t}) = \alpha + \beta_1 * (\text{v_Yr_Sold}==2007) + \beta_2 * (\text{v_Yr_Sold}==2008)$
7. Choose your own adventure: Pick any five variables from the dataset that you think will generate good R^2 . Use them in a regression of $\log(\text{Sale Price}_{i,t})$

- Tip: You can transform/create these five variables however you want, even if it creates extra variables. For example: I'd count Model 6 above as only using one variable: `v_Yr_Sold`.
- I got an R2 of 0.877 with just "5" variables. How close can you get? One student in five years has beat that.

Bonus formatting trick: Instead of reporting all regressions separately, report all seven regressions in a *single* table using `summary_col`.

```
In [21]: df = (
    df
    .assign(
        l_Lot_Area = np.log(df['v_Lot_Area']),
        l_SalePrice = np.log(df['v_SalePrice']),
        Yr_Sold = df['v_Yr_Sold'],
        const = 1
    )
)

df['Qual_bins'] = pd.cut(
    df['v_Overall_Qual'],
    bins=[0, 4, 6, 8, 10],
    labels=['Low', 'Medium', 'High', 'Very High']
)
```

```
In [22]: import statsmodels.formula.api as smf
from statsmodels.iolib.summary2 import summary_col

df['log_SalePrice'] = np.log(df['v_SalePrice'])
df['log_Lot_Area'] = np.log(df['v_Lot_Area'])

reg1 = smf.ols('v_SalePrice ~ v_Lot_Area', data=df).fit()
reg2 = smf.ols('v_SalePrice ~ log_Lot_Area', data=df).fit()
reg3 = smf.ols('log_SalePrice ~ v_Lot_Area', data=df).fit()
reg4 = smf.ols('log_SalePrice ~ log_Lot_Area', data=df).fit()
reg5 = smf.ols('log_SalePrice ~ v_Yr_Sold', data=df).fit()
reg6 = smf.ols('log_SalePrice ~ C(v_Yr_Sold)', data=df).fit()
reg7 = smf.ols('np.log(v_SalePrice) ~ v_Overall_Qual + v_Gr_Liv_Area + v_Garage_Car

summary_col([reg1, reg2, reg3, reg4, reg5, reg6, reg7],
            stars=True,
            model_names=['SP~Area', 'SP~log(Area)', 'log(SP)~Area', 'log(SP)~log(Ar
```

Out[22]:

	SP~Area	SP~log(Area)	log(SP)~Area	log(SP)~log(Area)	log
Intercept	154789.5502***	-327915.8023***	11.8941***	9.4051***	
	(2911.5906)	(30221.3471)	(0.0146)	(0.1511)	
v_Lot_Area	2.6489***		0.0000***		
	(0.2252)		(0.0000)		
log_Lot_Area		56028.1700***		0.2883***	
		(3315.1392)		(0.0166)	
v_Yr_Sold					
C(v_Yr_Sold) [T.2007]					
C(v_Yr_Sold) [T.2008]					
v_Overall_Qual					
v_Gr_Liv_Area					
v_Garage_Cars					
v_Total_Bsmt_SF					
v_1st_Flr_SF					
R-squared	0.0666	0.1284	0.0646	0.1350	
R-squared Adj.	0.0661	0.1279	0.0641	0.1345	

Standard errors in parentheses.

* p<.1, ** p<.05, ***p<.01

1. Sale Price $_{i,t}$ = 154789.55 + 2.6489 * v_Lot_Area
2. Sale Price $_{i,t}$ = -327915.80 + 56028.17 * log(v_Lot_Area)
3. log(Sale Price $_{i,t}$) = 11.8941 + 0.00 * v_Lot_Area
4. log(Sale Price $_{i,t}$) = 9.4051 + 0.2883 * log(v_Lot_Area)

5. $\log(\text{Sale Price}_{i,t}) = 22.2932 - 0.0051 * v_Yr_Sold$
6. $\log(\text{Sale Price}_{i,t}) = 12.0229 + 0.0256 * (v_Yr_Sold == 2007) - 0.0103 * (v_Overall_Qual + 0.0002 * v_Gr_Liv_v_1st_Flr_SF)$
7. $\log(\text{Sale Price}_{i,t}) = 10.5440 + 0.1383 * v_Overall_Qual + 0.0002 * v_Gr_Liv_v_1st_Flr_SF$

Part 3: Regression interpretation

Insert cells as needed below to answer these questions. Note that i is indexing a given house, and t indexes the year of sale.

1. If you didn't use the `summary_col` trick, list β_1 for Models 1-6 to make it easier on your graders.
2. Interpret β_1 in Model 2.
3. Interpret β_1 in Model 3.
 - HINT: You might need to print out more decimal places. Show at least 2 non-zero digits.
4. Of models 1-4, which do you think best explains the data and why?
5. Interpret β_1 in Model 5
6. Interpret α in Model 6
7. Interpret β_1 in Model 6
8. Why is the R2 of Model 6 higher than the R2 of Model 5?
9. What variables did you include in Model 7?
10. What is the R2 of your Model 7?
11. Speculate (not graded): Could you use the specification of Model 6 in a predictive regression?
12. Speculate (not graded): Could you use the specification of Model 5 in a predictive regression?

```
In [25]: print("Model 1  $\beta_1$ :", reg1.params[1])
print("Model 2  $\beta_1$ :", reg2.params[1])
print("Model 3  $\beta_1$ :", reg3.params[1])
print("Model 4  $\beta_1$ :", reg4.params[1])
print("Model 5  $\beta_1$ :", reg5.params[1])
print("Model 6  $\beta_1$ :", reg6.params['C(v_Yr_Sold)[T.2007]'])
print("Model 6  $\beta_2$ :", reg6.params['C(v_Yr_Sold)[T.2008]'])
print("Model 7  $\beta_1 - \beta_5$ :")
print(reg7.params[1:6])
```

```

Model 1  $\beta_1$ : 2.648935000718191
Model 2  $\beta_1$ : 56028.16996046537
Model 3  $\beta_1$ : 1.3092338465836504e-05
Model 4  $\beta_1$ : 0.28826331962293017
Model 5  $\beta_1$ : -0.005114348195977281
Model 6  $\beta_1$ : 0.025590319971647263
Model 6  $\beta_2$ : -0.010281565074487964
Model 7  $\beta_1$  -  $\beta_5$ :
v_Overall_Qual      0.138264
v_Gr_Liv_Area       0.000191
v_Garage_Cars       0.105987
v_Total_Bsmt_SF     0.000098
v_1st_Flr_SF        0.000052
dtype: float64

```

```

C:\Users\lenovo\AppData\Local\Temp\ipykernel_52060\2003613305.py:1: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
  print("Model 1  $\beta_1$ :", reg1.params[1])
C:\Users\lenovo\AppData\Local\Temp\ipykernel_52060\2003613305.py:2: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
  print("Model 2  $\beta_1$ :", reg2.params[1])
C:\Users\lenovo\AppData\Local\Temp\ipykernel_52060\2003613305.py:3: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
  print("Model 3  $\beta_1$ :", reg3.params[1])
C:\Users\lenovo\AppData\Local\Temp\ipykernel_52060\2003613305.py:4: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
  print("Model 4  $\beta_1$ :", reg4.params[1])
C:\Users\lenovo\AppData\Local\Temp\ipykernel_52060\2003613305.py:5: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
  print("Model 5  $\beta_1$ :", reg5.params[1])

```

2. If $v_{\text{Lot_Area}}$ goes up by 1%, the house price increases by about \$560.
3. If $v_{\text{Lot_Area}}$ increases by 1 unit, the house price goes up by about 0.00139%.
4. I think Model 4 is the best because it uses the log form and has the highest R-squared among the four models.
5. It means that for every year increase, the log of SalePrice decreases by about 0.0051.

6. α represents the log average of house sale prices in 2006 (with 2006 as the baseline year), which means the average sale price in 2006 was approximately \$165,500.
7. β_1 represents the average percentage change in house prices in 2007 compared to 2006. That means house prices in 2007 were on average about 2.56% higher than in 2006.
8. Model 6 has a higher R^2 because it considers the effects of 2007 and 2008 separately, making it more detailed. Model 5 only treats the year as a continuous variable and doesn't capture the specific differences between years. So Model 6 explains the changes in housing prices better.
9. In model 7, I included: $v_Overall_Qual$, $v_Gr_Liv_Area$, v_Garage_Cars , $v_Total_Bsmt_SF$, $v_1st_Flr_SF$
10. The R^2 in Model 7 is 0.8024
11. I think Model 6 is not suitable for prediction because it uses two dummy variables for the years 2007 and 2008.
12. I think Model 5 is not suitable for prediction because its R-squared is 0.0014 and its adjusted R-squared is negative. This means the model has no predictive power.

In []: